

Neural.Orb

Exploration of Deep-Q Networks in a Physical Space

a project by

Tabor Henderson

thetabor@gmail.com

git://thetabor

tech stack

♥ Python 3.6

Keras

Theano

TensorFlow

SKImage

Jupyter Notebook

pygame.camera

Python.Swarms

kulka

PyBluez

tqdm

Lubuntu

robots

Sphero Sprk

special thanks

Frank, Adam,

Steve and Brent!

Wim, for all the advice

♥ Embur

and my family

background

DQNs

Deep-Quality Networks

Neural Network Training Technique

DeepMind demonstrated their early successes with DQNs in 2013. They accomplished "superhuman" level performance on many Atari games, and then developed the AlphaGo network which recently beat the No. 1 ranked player in the world at a Go summit.

Training

Rewards provided by environment

Rewards give the agent feedback about its environment. By learning to predict these rewards, the DQN agent can discover strategies without direct guidance.

project goals

Simulation

Demonstrate DQN learning

Proof of Concept

Using the Python.Swarms cellular automata engine, I built methods for training and deploy supervised and DQN learners. The DQNs take much long to train. However, they can respond to changes in their environment.

Navigation

Deploying to a physical space

Application

The Sphero robotic ball provides a perfect platform for a DQN experiment. By scoring the Sphero's distance from the center of an image, I'm training it to chase that spot. Results are positive, but training is very slow.

Cooperation

Training complex behavior

Potential next step

Finding ways to train cooperative behavior opens many possibilities. This requires solving a compounding of all the challenges already navigated for a single agent.

references

1

Deep Reinforcement Learning: Pong from Pixels

Andrej Karpathy

blog

2

Keras plays catch, a single file Reinforcement Learning example

Eder Santana

github.io

3

Demystifying Deep Reinforcement Learning

Tambet Matiisen

University of Tartu

4

Playing Atari with Deep Reinforcement Learning

Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, Martin Riedmiller

NIPS Deep Learning Workshop 2013