



# Neural.Swarms and Neural.Orb

Tabor C. Henderson  
Galvanize Denver DSi

## Deep Quality Networks

- Deep-Q Networks apply reinforcement learning to neural networks.
- The network predicts Q-values for each action the network is allowed.
- A Q-value represents quality of a state, or the expected sum of rewards as we play the game from that state.
- Our agents (almost) always select the max Q-value we predict.
- DQN's were deployed to great effect by DeepMind, most recently in their Go-playing AI, AlphaGo.

## Objectives

- Simulation: demonstrate DQNs
- Navigation: utilize DQNs in a physical space
- Coordination: train multiple DQN-based agents to cooperate

## Simulation

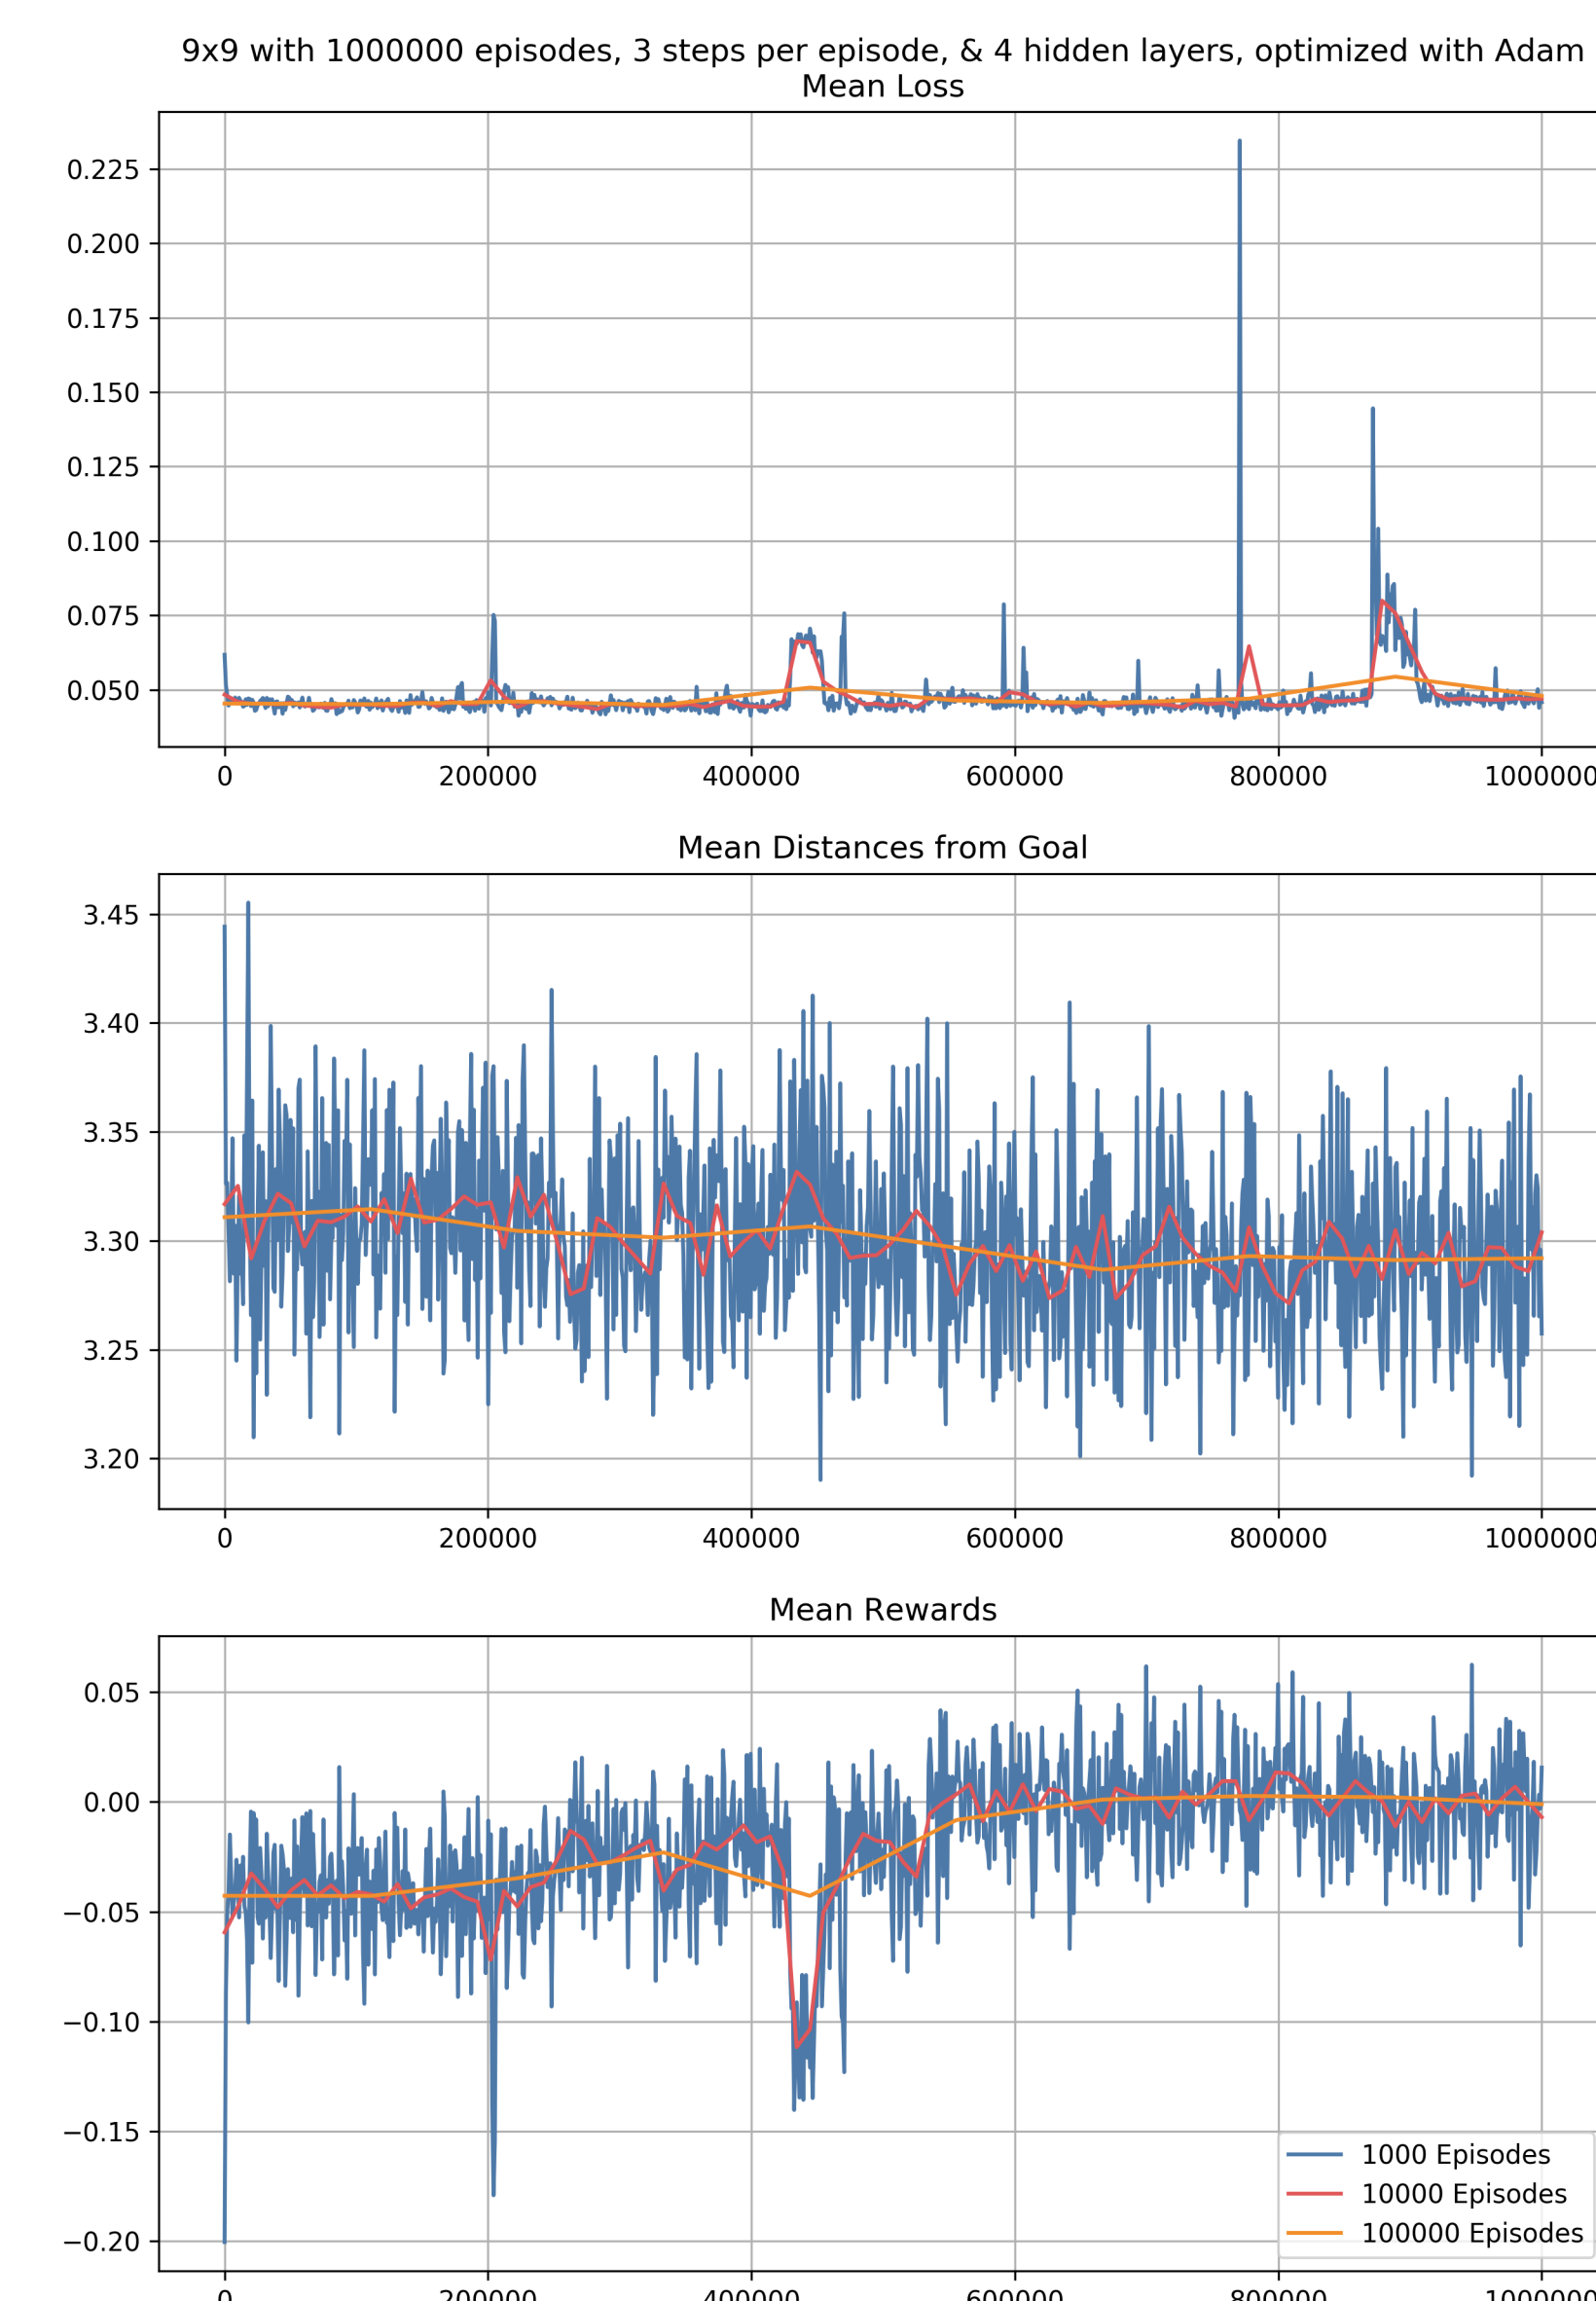
I developed the Neural.Swarms simulation engine in which I can implement either supervised or reinforcement learning.





## Measures

Measuring the DQN's learning was difficult, as they require a massive amount of training episodes (similar to epochs in supervised learning). I used three different measures of network learning

- Loss: mean squared error, commonly used for many machine learning algorithms
- Mean distance to target: measured over one or more game episodes, this provides feedback on actual performance in game
- Mean reward: measured in the same intervals as mean distance, mean reward tells us the same basic information as mean distance, and their inverse relationships confirm that the scoring system works.



## Computer Vision



## Deploying to Sphero

- Deploying a DQN to Sphero required a computer vision interface, and controlled training environment. SciKit-Image provides excellent filters and segmentation modules, and pygame.camera allows access to webcams.
- Bluetooth instability frequently interrupted training runs, and at one point the host machine overheated and crashed
- The Sphero DQN couldn't utilize the experience replay technique, which has been shown to accelerate learning
- Each game step took 4-5 seconds, meaning performing the million steps required for good performance would take roughly 10 days of continuous training.

## Conclusion

While I can't live my fantasy of terrorizing friend's dogs with swarms of robot balls, I can control one, albeit inaccurately, with a neural network. The training program works, but simply doesn't run fast enough for me to have produced a viable model in the allotted time frame. However, a viable model is possible given a few improvements:

- Implementation of experience replay for the physical model (required too much disk space)
- A simulation engine similar enough to the physical problem that we could pre-train the model virtually, then deploy to the physical space
- Utilize all data from the Sphero's onboard sensors- I was only able to poll data from the locator which provides (x, y) coordinates but not acceleration or angle.

## References

1. Andrej Karpathy
2. Eder Santana
3. Charles Isbell, Michael Littman, and Chris Pryboy via Udacity
4. Tabet Matiisen